

Reading and importing the data

```
1 import nltk
2
3 # Download the stopwords resource
4 nltk.download('stopwords')
5
6 # Download the tokenizer resource
7 nltk.download('punkt')
8
9 import numpy as np
10 import pandas as pd
11 import matplotlib.pyplot as plt
12 import seaborn as sns

[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data] Package stopwords is already up-to-date!
[nltk_data] Downloading package punkt to /root/nltk_data...
[nltk_data] Package punkt is already up-to-date!
```

```
1 data=pd.read_csv("/content/Twitter_Data.csv")
```

```
1 data.info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 162969 entries, 0 to 162979
Data columns (total 3 columns):
#   Column          Non-Null Count  Dtype
---  ---
0   clean_text      162969 non-null object
1   category        162969 non-null object
2   sentence_length 162969 non-null int64
dtypes: int64(1), object(2)
memory usage: 5.0+ MB
```

```
1 data.describe()
```

	category
count	162973.000000
mean	0.225436
std	0.781279
min	-1.000000
25%	0.000000
50%	0.000000
75%	1.000000
max	1.000000

Changing our dependent variable to categorical. (0 to "Neutral",-1 to "Negative", 1 to "Positive")

```
1 # Map numeric values to categorical labels
2 category_mapping = {0: 'Neutral', -1: 'Negative', 1: 'Positive'}
3 data['category'] = data['category'].map(category_mapping)
4
```

Doing Missing value analysis and drop all null/missing values

```
1 # Check for missing values
2 missing_values = data.isnull().sum()
3
4 # Drop rows with missing values
5 data.dropna(inplace=True)
6
```

Dummy Variables

```
1 # Create dummy variables for the dependent variable
2 data = pd.get_dummies(data, columns=['category'], prefix='', prefix_sep='')
3
```

Doing text cleaning. (remove every symbol except alphanumeric, transform all words to lower case, and remove punctuation and stopwords)

```
1 # Remove symbols, transform to lowercase, remove punctuation, and stopwords
2 def clean_text(text):
3     text = re.sub(r'[^\w\s]', '', text) # Remove symbols
4     text = text.lower() # Transform to lowercase
5     text = re.sub(r'\s+', ' ', text) # Remove extra spaces
6     tokens = word_tokenize(text) # Tokenize the text
7     tokens = [word for word in tokens if word not in stopwords.words('english')]
8     return ' '.join(tokens)
9
10 data['cleaned_text'] = data['clean_text'].apply(clean_text)
```

Creating a new column and finding the length of each sentence (how many words they contain)

```
1 data['sentence_length'] = data['clean_text'].apply(lambda x: len(x.split()))
2
```

Splitting data into dependent(X) and independent(y) dataframe

```
1 X = data['clean_text']
2 y = data['category']
```

3

Do operations on text data

```

1 from tensorflow.keras.preprocessing.text import Tokenizer
2 from tensorflow.keras.preprocessing.sequence import pad_sequences
3 from sklearn.model_selection import train_test_split
4
5 # Create Tokenizer
6 tokenizer = Tokenizer()
7 tokenizer.fit_on_texts(X)
8 X_encoded = tokenizer.texts_to_sequences(X)
9 X_padded = pad_sequences(X_encoded, padding='pre')
10

```

Build an LSTM Model:

```

1 from tensorflow.keras.models import Sequential
2 from tensorflow.keras.layers import Embedding, LSTM, Dense
3
4 # Define the model architecture
5 model = Sequential()
6 model.add(Embedding(input_dim=len(tokenizer.word_index)+1, output_dim=100, input_1
7 model.add(LSTM(100))
8 model.add(Dense(3, activation='softmax'))
9
10 # Compile the model
11 model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
12

```

Split data into train and test

```

1 X_train, X_test, y_train, y_test = train_test_split(X_padded, y, test_size=0.2, ra

```

Training new model

```

1 history = model.fit(X_train, pd.get_dummies(y_train), epochs=10, batch_size=32, va
2

```

```

Epoch 1/10
3260/3260 [=====] - 985s 301ms/step - loss: 0.3759 - a
Epoch 2/10
3260/3260 [=====] - 1013s 311ms/step - loss: 0.2170 - i
Epoch 3/10
3260/3260 [=====] - 1005s 308ms/step - loss: 0.1396 - i
Epoch 4/10
3260/3260 [=====] - 1027s 315ms/step - loss: 0.0854 - i
Epoch 5/10
3260/3260 [=====] - 1031s 316ms/step - loss: 0.0533 - i
Epoch 6/10
3260/3260 [=====] - 982s 301ms/step - loss: 0.0329 - a
Epoch 7/10
3260/3260 [=====] - 992s 304ms/step - loss: 0.0220 - a

```

```

Epoch 8/10
3260/3260 [=====] - 1004s 308ms/step - loss: 0.0157 - i
Epoch 9/10
3260/3260 [=====] - 996s 306ms/step - loss: 0.0121 - a
Epoch 10/10
3260/3260 [=====] - 988s 303ms/step - loss: 0.0088 - a

```

Normalize the Predictions

```

1 y_pred = model.predict(X_test)
2 y_pred_normalized = np.where(y_pred >= 0.5, 1, 0)
3

1019/1019 [=====] - 15s 15ms/step

1 y_pred_labels = np.argmax(y_pred, axis=1)
2

```

Measure performance metrics and accuracy

```

1 from sklearn.metrics import accuracy_score
2 # Convert y_test to a numpy array
3 y_test_array = np.array(y_test)
4 # Reshape y_test if needed
5 if len(y_test_array.shape) == 1:
6     y_test_array = y_test_array.reshape(-1, 1)
7
8 y_test_labels = np.argmax(y_test_array, axis=1)
9 accuracy = accuracy_score(y_test_labels, y_pred_labels)
10 print("Accuracy:", accuracy)
11

```

Accuracy: 0.21642020003681658

Print Classification report

```

1 from sklearn.metrics import classification_report
2
3 classification_rep = classification_report(y_test_labels, y_pred_labels)
4 print("Classification Report:")
5 print(classification_rep)
6

```

```

Classification Report:
              precision    recall  f1-score   support

     0           1.00        0.22        0.36        32594
     1           0.00        0.00        0.00           0
     2           0.00        0.00        0.00           0

 accuracy                   0.22        32594
 macro avg              0.33        0.07        0.12        32594
 weighted avg           1.00        0.22        0.36        32594

```

```
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344
_warn_prf(average, modifier, msg_start, len(result))
```

